

Mentoring Operating System (MentOS)

Process management

Alessandro Danese, Luigi Capogrosso

University of Verona
alessandro.danese@univr.it,
luigi.capogrosso@studenti.univr.it

Version 1.0.2



Table of Contents

- 1 Process descriptor
 - Process identifier
 - State of a process
 - Time accounting
 - Context of a process
- 2 Scheduler
 - Data structures
 - Scheduling algorithms
 - Context switch



Process descriptor



Process descriptor

task_struct is the data structure used by the Kernel to represent a process¹.

```
struct task_struct {  
    pid_t pid;                // the process identifier  
    unsigned long state;      // the current process's state  
    struct task_struct *parent; // pointer to parent process  
    struct list_head children;  // list of children process  
    struct list_head siblings;  // list of siblings process  
    struct mm_struct *mm;       // memory descriptor  
    struct sched_entity se;     // time accounting  
    struct thread_struct thread; // context of process  
}
```

N.B. The memory descriptor of a process is only here reported.
It will be illustrated in the section Memory Management.

¹In Linux, it is quite big, 1.7KB on 32-bit machine (include/linux/sched.h)



Process identifier

Process Identifier (PID) is numeric value identifying a process. When a new process is created a new PID is generated by summing 1 to the last assigned PID.

In Linux, the maximum value for a PID is 32768. When the PID maximum value is reached, the last assigned PID is reset to 0 before searching for a new PID.

The macro `RESERVED_PID` (usually set to 300) is defined to reserved PIDs to system processes and daemons, namely processes proving a service (e.g. a web server). All user's processes have PID greater than `RESERVED_PID`.



State of a process

Process state is a numeric value describing the current state of the process. A process can be in one of the following state:

- **TASK_RUNNING**: either the process is currently in execution, or it has all the resources to be executed except the CPU.
- **TASK_INTERRUPTIBLE**: the process is blocked (sleep), waiting for some condition to run. When this condition exists, the kernel sets the process's state to **TASK_RUNNING**. The process also awakes and becomes runnable if it receives a signal.
- **TASK_UNINTERRUPTIBLE**: this state is identical to **TASK_INTERRUPTIBLE** except that it does not wake up and become runnable if it receives a signal. This is used in situations where the process must wait without interruption.

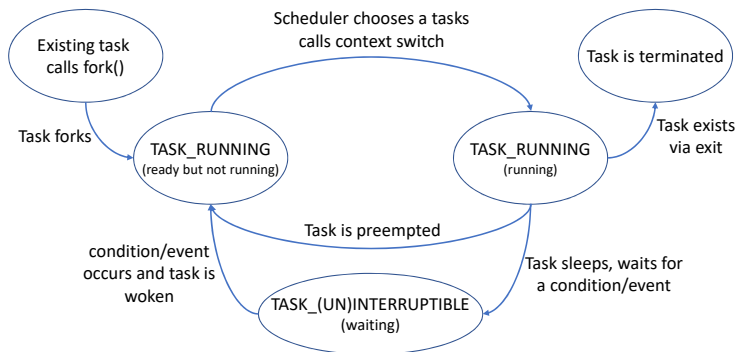


Process state (continue)

- **TASK_STOPPED**: process execution has stopped; the task is not running nor is it eligible to run.
- **EXIT_ZOMBIE**: Process execution is terminated, but the parent process has not yet issued a `wait4(0)` or `waitpid()` system call to return information about the dead process.
- **EXIT_DIED**: The final state: the process is being removed by the system because the parent process has just issued a `wait4()` or `waitpid()` system call for it.



State of a process



Flow chart of process states



Relationships among processes


Processes created by a program have a parent/child relationship. When a process creates multiple children, these children have sibling relationships.

Fields of `task_struct` describing the relationships among processes:

- **parent:** pointer to the process's parent;
- **children:** The head of the list containing all children created by the process.
- **sibling:** The head of the list containing all children created by the process's parent.



Relationships among processes



Parenthood relationships among four processes.



Time accounting

The filed struct sched_entity **se** of a task_struct reports the priority and the execution times of a process.

```
struct sched_entity {  
    int      prio;           // priority  
    time_t   start_runtime;  // start execution time  
    time_t   exec_start;     // last context switch time  
    time_t   sum_exec_runtime; // overall execution time  
    time_t   vruntime;       // weighted execution time  
}
```



Time accounting

- **prio** defines the execution priority of a process. It has a value in the range $[100, 139]$, where 100 means the highest priority, and 139 means the lowest priority.

By default, the priority of a new generated process is 120.

A process can increment/decrement its *prio* value by using the system call *nice(inc)*, which takes as input parameter a value in the range $[-20, 19]$.

Examples:

nice(1) (increment *prio* value of calling process by 1 unit)

nice(-5) (decrement *prio* value of calling process by -5 units)



Time accounting

- **start_runtime:**

The system execution time reporting when the process was first executed in the CPU.

- **exec_start:**

The system execution time reporting when the process was last executed in the CPU.

- **sum_exec_runtime:**

The overall execution time spent by the process in CPU.

- **vruntime:**

The virtual runtime, namely the weighted overall execution time spent by the process in CPU (see CFS).



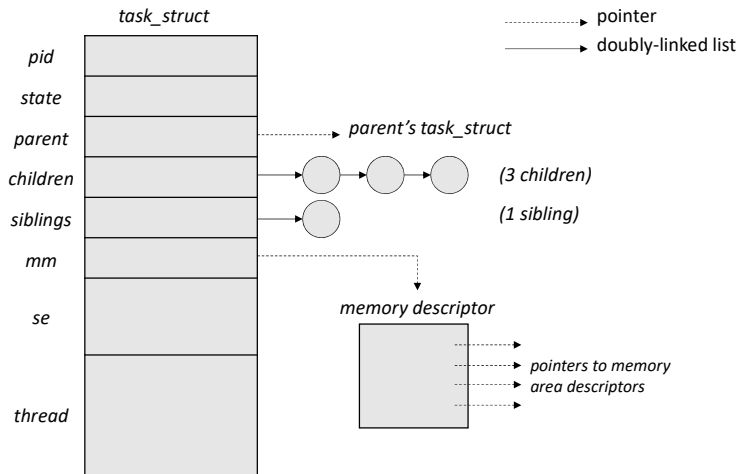
Context of a process

The filed struct `thread_struct` **thread** of a `task_struct` reports the context of a process whenever it is being switched out.

```
struct thread_struct {
    uint32_t ebp;           // base pointer register
    uint32_t esp;           // stack pointer register
    uint32_t ebx;           // base register
    uint32_t edx;           // data register
    uint32_t ecx;           // counter
    uint32_t eax;           // accumulator register
    uint32_t eip;           // Instruction Pointer Register
    uint32_t eflags;        // flag register
    bool_t fpu_enabled;     // is FPU enabled?
    savefpu fpu_register;   // FPU context
}
```



task_struct memory representation



Scheduler



Scheduler data structures

The **runqueue** data structure is the most important data structure of the scheduler. It collects all system processes in running state.

```
struct runqueue {  
    unsigned long nr_running; // number of processes in running state  
    struct task_struct *curr;  // pointer to current running process  
    struct list_head queue;    // list of processes in running state  
}
```

Pay attention!

queue is the *head* of a circular, doubly-linked list collecting all system processes in running state. Consequently, a field **run_list** of type *struct list_head* is added in the *struct task_struct*.

(see slides fundamental concepts for more details).



Scheduler execution flow

The scheduler is called after the handle of an interrupt/exception. In detail, the following operations are performed by the scheduler:

- 1 update the time accounting variables of the current process;
- 2 try to wake up a waiting process. Whether a waiting condition is met, a process is woken by setting its state to running, and inserting it into the runqueue (topic not faced in current slides);
- 3 run scheduling algorithm to pick the next process to be executed by CPU from the runqueue;
- 4 perform context switch.



Scheduler - select next process

pick_next_task is the function called by scheduler to get the next process to execute. According to the implemented scheduling algorithm, the next process can be differently chosen.

MentOs provides the following tree algorithms:

- **RR** - Round-Robin.
- **Priority** - Highest Priority First.
- **CFS** - Completely Fair Scheduler.

Pay attention!

In the following algorithms, we use the doubly-linked list defined in Linux Kernel, to collect all processes in running state.



Scheduler - select next process (Round-Robin)

Round Robin is a CPU scheduling algorithm where a fixed time slice is assigned to each system process, in a cyclic way. It is simple, preemptive, easy to implement, and starvation-free.



Scheduler - select next process (Round-Robin)

Algorithm 1 pseudocode of Round-Robin algorithm.

Require: Current process c , List of processes L

Ensure: Next process n

- 1: $nNode = next(c)$
 - 2: **if** $isHead(L, nNode)$ **then**
 - 3: $nNode = next(nNode)$
 - 4: **end if**
 - 5: $n = entry(nNode)$
-



Scheduler - select next process (Round-Robin)

C Implementation of Round-Robin algorithm.

```
struct task_struct * pick_next_task(struct runqueue *runqueue) {  
    // nNode = next(c)  
    struct list_head *nNode = runqueue->curr->run_list.next;  
    // if isHead(L, nNode)  
    if (nNode == &runqueue->queue)  
        nNode = nNode->next;  
    // n = entry(nNode)  
    task_struct *next;  
    next = list_entry(nNode, struct task_struct, run_list);  
    return next;  
}
```



Scheduler - example of execution (Round-Robin)

First iteration:

- `current_process = init`
- `pick_next_task()` returns `shell`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Second iteration:

- `current_process = shell`
- `pick_next_task()` returns `task1`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Third iteration:

- `current_process = task1`
- `pick_next_task()` returns `task2`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Fourth iteration:

- `current_process = task2`
- `pick_next_task()` returns `task3`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Fifth iteration:

- `current_process = task3`
- `pick_next_task()` returns `task4`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Sixth iteration:

- `current_process = task4`
- `pick_next_task()` returns `init`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - example of execution (Round-Robin)

Seventh iteration:

- `current_process = init`
- `pick_next_task()` returns `shell`

runqueue

init	shell	task1	task2	task3	task4
------	-------	-------	-------	-------	-------



Scheduler - select next process (Highest Priority First)

Round robin scheduling assumes that all processes are equally important. This generally is untrue. We would sometimes like to see long CPU-intensive (non-interactive) processes get a lower priority than interactive processes.

In addition, different users may have different status. A system administrator's processes may rank above those of a student's.

These goals led to the introduction of the *Priority* scheduling algorithm.



Scheduler - select next process (Highest Priority First)

Each process has a static priority. Smaller is the number, higher is the priority of the process.

The scheduler simply picks the highest priority process to run. A process is preempted whenever a higher priority process is available in the run queue.

Advantage: priority scheduling provides a good mechanism where the relative importance of each process may be precisely defined.

Disadvantage: If high priority processes use up a lot of CPU time, lower priority processes may starve and be postponed indefinitely, leading to starvation.



Scheduler - select next process (Highest Priority First)

Algorithm 2 pseudocode of Highest Priority First.

Require: Current process c , List of processes L

Ensure: Next process n

```
1:  $n = c$ 
2: for all  $INode \in L$  do
3:    $t = \text{entry}(INode)$ 
4:   if  $\text{priority}(t) < \text{priority}(n)$  then
5:      $n = t$ 
6:   end if
7: end for
```

The C implementation of this scheduling algorithm is given to the student.



Scheduler - example of execution (Highest Priority First)

First block of iteration:

- `current_process = task2`
- `pick_next_task()` returns `task2` until no process with an higher priority is present in the system.

runqueue

init prio = 120	shell prio = 120	task1 prio = 110	task2 prio = 105
--------------------	---------------------	---------------------	---------------------



Scheduler - example of execution (Highest Priority First)

Second block iteration:

- `current_process = task1`
- `pick_next_task()` returns `task1` until no process with an higher priority is present in the system.

runqueue

init prio = 120	shell prio = 120	task1 prio = 110	
--------------------	---------------------	---------------------	--



Scheduler - example of execution (Highest Priority First)

Third block of iteration:

- `current_process = task3`
- `pick_next_task()` returns `task3` until no process with an higher priority is present in the system.

runqueue

init prio = 120	shell prio = 120	task1 prio = 110	task3 prio = 105
--------------------	---------------------	---------------------	---------------------

How much time do init and shell have to wait to get the CPU?



Scheduler - select next process (Completely Fair Scheduler)

Completely Fair Scheduler (CFS) aims to prevent starvation by assigning the CPU fairly to all system processes.

Let consider an example to illustrate the goal of CPF. If there are two tasks A and B, which have a same “weight”, the portion of available CPU time given to each task is 50%.

However, if the “weight” of task increases on CPU by 10%, then task A's portion of the CPU is 55%, meanwhile task B's portion of the CPU becomes 45%.



Scheduler - select next process (Completely Fair Scheduler)

CFS's idea: let use the priority of each process to “weight” its overall execution time (virtual runtime).

Processes having a low priority have a virtual runtime increasing faster than processes with a higher priority. Scheduler always picks the process with the lowest virtual execution time!



Scheduler - select next process (Completely Fair Scheduler)

Scheduler needs to know the weight of the task to estimate its CPU time's portion. Hence, the priority number has to be mapped to such a weight; this is done in the array `prio_to_weight`:

```
static const int prio_to_weight[] = {  
    /* 100 */    88761, 71755, 56483, 46273, 36291,  
    /* 105 */    29154, 23254, 18705, 14949, 11916,  
    /* 110 */    9548, 7620, 6100, 4904, 3906,  
    /* 115 */    3121, 2501, 1991, 1586, 1277,  
    /* 120 */    1024, 820, 655, 526, 423,  
    /* 125 */    335, 272, 215, 172, 137,  
    /* 130 */    110, 87, 70, 56, 45,  
    /* 135 */    36, 29, 23, 18, 15  
};
```



Scheduler - select next process (Completely Fair Scheduler)

A priority number of 120, which is the priority of a normal task, is mapped to a weight of 1024.

Note that the ratio of two successive entries in the array is almost 1.25. This number is chosen such that:

- if the priority of a task is reduced by one, then it gets 10% higher share of the available CPU time.
- if the priority of a task is increased by one, then it gets 10% lower share of the available CPU time.



Scheduler - select next process (Completely Fair Scheduler)

Given the array `prio_to_weight` we can update the virtual runtime of a process `p`, namely its weighted overall execution by using the formula:

$$\text{vruntime} += \text{delta_exec} * (\text{NICE_0_LOAD} / \text{weight}(p))$$

where:

- `vruntime` is the virtual run time of the process;
- `delta_exec` is the last amount of time spent by `p` in the CPU;
- `NICE_0_LOAD` is the weight of a task with normal priority (1024);
- `weight(p)` is the weight of `p` defined by the array `prio_to_weight`.



Scheduler - select next process (Completely Fair Scheduler)

Algorithm 3 pseudocode of Completely Fair Scheduler.

Require: Current process c , List of processes L


Ensure: Next process n

```
1: updateVirtualRuntime( $c$ )
2:  $n = c$ 
3: for all  $I\text{Node} \in L$  do
4:    $t = \text{entry}(I\text{Node})$ 
5:   if  $\text{virtualRuntime}(t) < \text{virtualRuntime}(n)$  then
6:      $n = t$ 
7:   end if
8: end for
```

The C implementation of this scheduling algorithm is given to the student.



Scheduler - context switch

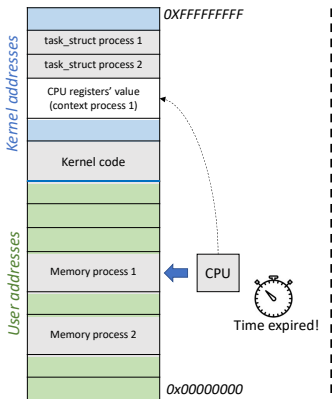


The CPU performs a context switch to change the process executed by CPU.

The following example shows the steps performed by the operating system to save the current process's state (*process 1*), and then resume the execution of a previously stopped process (*process 2*).



Scheduler - context switch

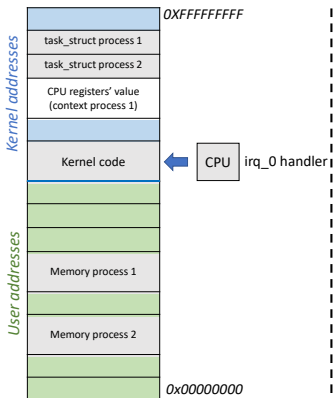


(1) Time expired! It is time to give back control of CPU to kernel. Timer device rises the signal *INTR* and present 0 in *irq* line.

When *INTR* is risen, the CPU moves from Ring 3 (user mode) to Ring 0 (kernel mode). After the CPU privilege level change, the values of CPU registers are pushed in the Kernel's stack.



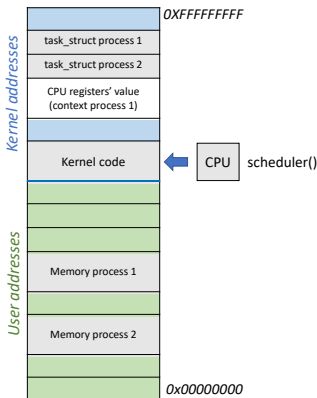
Scheduler - context switch



(2) CPU starts executing *irq_0* interrupt handler to handle the hardware interrupt 0 risen by Timer.



Scheduler - context switch

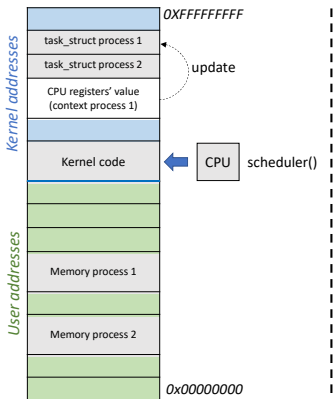


(3) The scheduler is then called to update the time accounting variables of the interrupted process, and pick the next process to run.

In this example, the scheduler picks the process 2 as the next one.



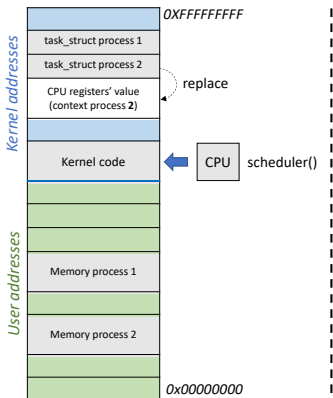
Scheduler - context switch



(4) Kernel updates the `thread_struct` structure of the `task_struct` of the process 1 in order to save its context.



Scheduler - context switch




(4) Kernel updates the `thread_struct` structure of the `task_struct` of the process 1 in order to save its context.

(5) Kernel replaces the context of process 1 with the context of process 2 in its stack memory.




Scheduler - context switch



(6) Kernel moves the values from its stack to CPU's registers and runs an *iret* assembly instruction, which changes the CPU privilege level from Ring 0 (kernel mode) to Ring 3 (user mode).



Scheduler - context switch



(7) The context of the process 2 is in the CPU's registers finally. The CPU can keep on executing the code of the process 2 in user mode until the next context switch.

